

Results — SCAFFOLD, no data recorded yet

Please use CSS style instead of skinparam padding

NO EVALUATION HAS BEEN RUN
Every value below is a placeholder.
Run backend/eval/run_eval.py and fill them in.

Scorecard — one row per metric

Quiz key accuracy

score: __
band: __
n cases: __
verdict: __

Tutor faithfulness

score: __
band: __
n cases: __
verdict: __

Retrieval precision@k

score: __
band: __
k = 5 (lecture) / 6 (course)
verdict: __

Synthesis quality (LLM judge)

score: __
band: __
n cases: __
verdict: __

These four metrics are the ones the existing harness actually computes (see F18). Do not add a fifth row without adding a scorer for it — a metric with no code behind it is the same class of error as the README claims audited in the honesty ledger.

Run metadata

commit: _____
pipeline: Live / Fake
model(s): __
date: _____
overall: PASS / FAIL

Claims this figure can support once filled:
pipeline output quality against curated golden sets, measured reproducibly at a known commit.

Claims it can NEVER support, however good the numbers:
that students learned more. There is no controlled study, no comparison group, and `learning\_events` retention defaults to 0 (never archive) with no analysis pipeline behind it. Say so explicitly rather than letting a reader infer a learning-outcomes claim from a quality metric.

Placeholder figure, tracked deliberately so the gap is visible rather than forgotten.
Replace before submission or remove the figure and its reference from Chapter 6.